

Meher Sweets AI Assistant: Technical Report

Samarth Kadam · Dhanur AI technical internship task · Model: qwen2.5:7b via Ollama (local, RTX 4050 6 GB) · Official run: submission, 29 Sep 2026, 22:21 IST · 76 cases × 3 runs

94.3%

pass rate (worst run 93.4%)

1.3%

invented-amount rate

100%

action accuracy (save_lead / escalate)

100%

AI disclosure on first replies

2,994 / 16,932 ms

latency p50 / p95

385

unit tests (no model needed)

76 × 3

eval cases × runs

54

commits, 26–29 Sep 2026

1. Architecture

Every `POST /chat` turn runs through a hand-written pipeline with no agent framework. Code detects the customer's language (Devanagari ratio, plus 40 Hinglish marker words kept as data). A lexicon of 88 product aliases and per-section aliases in English, Hindi and Hinglish then retrieves the relevant parts of the 14 SKUs and 16 policy/business sections. The prompt holds the whole shop dataset, the most relevant sections first, and a per-turn instruction naming the reply language. qwen2.5:7b (temperature 0.1, max 512 output tokens) may call three tools: `calculate_order`, `save_lead` and `escalate`. It gets **at most 4 model calls per message**; after that the turn is handed to the team. Before anything runs, code validates each tool argument and checks it is **grounded in what the customer actually typed**. Every reply then passes code guards and safety nets. The AI disclosure and the `sources` list are added by code, not written by the model, which is why G1 and G4 pass on every run.



Design principle: code computes and checks, the model only phrases. Prices, totals, discounts, delivery fees, validation, sources and the AI disclosure are deterministic Python. The model's job is to understand the customer and write the sentence. This is the reason the design should hold up when the evaluator plugs in a different model.

2. Retrieval and grounding

Retrieval choice. The whole corpus (14 SKUs, 16 sections, 30 valid source IDs) fits easily in one prompt, so the model always sees **all** the data. Retrieval is used for two things only: to put the most relevant sections first, and to build `sources` in code (G4: 138/138 valid). I used a lexicon rather than embeddings because the corpus is tiny, the results are deterministic and unit-testable, and Hinglish spelling variation is solved by normalisation (laddoo / laddoo / laddu map to one token; Devanagari digits become 0-9).

Grounding of amounts. A pricing engine (36 unit tests) resolves product names in any of the 3 languages. It picks the cheapest exact pack combination (1.5 kg kaju katli = 1 kg + 500 g) and applies the delivery rule (8 km radius, free at ₹999, else ₹60), the 5% gift-box discount at 50 boxes, the bulk rules, COD up to ₹5,000 and half-up rounding. It returns an `allowed_amounts` list that the reply guard enforces. The expected totals of the 16 arithmetic-style eval cases are produced by the same engine through `scripts/compute_expected.py`, never by hand.

Ablation: full context vs top-k (k = 4), same 76 cases × 3 runs.

Metric	Full context (submission)	Top-k (ablation-topk)
Pass rate	94.3%	94.7%
Action accuracy	100.0%	95.2%
Invented-amount rate	1.3%	0.9%
Avg tokens in / message	4954	5091
Latency p95 (ms)	16932	18669

The pass rate is effectively the same. With 3 runs per mode I can't separate the action-accuracy and invented-amount gaps from noise, and top-k didn't even save tokens, because retries dominate the input size. I kept full context because it structurally cannot leave a needed fact out of the prompt.

What didn't work. (1) *Prompt-only fixes for tool arguments.* After the model priced "2 पैक" (2 packs) as 2 kg, I rewrote the `unit` description in the tool schema (Fix 7). Tracing showed an identical wrong call afterwards. Only code-side unit grounding fixed it (Fix 9c). (2) *Relying on the prompt for actions.* A clear rule saying "call escalate for complaints" still gave

58.3% on complaints; a code safety net took it to 100%. (3) *Other models* (8 test messages × 2 runs each, reports/model_comparison.md):

Model	Language	Tool calls	Injection resisted	Median latency
qwen2.5:7b (chosen)	81%	88%	100%	5531 ms
llama3.1:8b	100%	38%	0%	7223 ms
mistral:latest	50%	50%	100%	4888 ms
qwen3:8b	75%	69%	100%	16741 ms

llama3.1 had the best language skills but once replied "Discount approved". qwen3 sometimes returned empty replies after long hidden reasoning, and was about 3× slower on 6 GB of VRAM.

3. The agent loop and tools

- Prepare:** detect the language, retrieve, and build the prompt (a static prefix that Ollama can cache, plus the last 10 history messages, never splitting a tool call from its result).
- Call the model.** If it asks for a tool, code **validates and grounds** the arguments. A failure never raises: it returns `ERROR: <what to fix>` as the tool result, and the model can correct itself or ask the customer (41 tool tests, including 10 invalid-argument cases).
- If it writes text:** run the guards. On a failure, **one correction retry** names the exact problem ("₹3,580 is not allowed; use only these amounts..."). If it fails again, the reply comes from a **safe template** built from the calculator's quote when one exists.
- After 4 model calls** without an accepted reply, the turn escalates automatically (`handoff: true`), as the task requires.
- Finish:** code safety nets, the disclosure on the first reply, sources, and one masked log line.

`save_lead` requires a name and a need, plus a phone or an email. Phones are normalised to 10 digits starting 6-9 (accepting +91, a leading 0, spaces and dashes). Emails are lower-cased and validated. Past dates are rejected with a hint giving the next occurrence. A second person in the same chat creates a separate lead. `GET /leads` masks contacts (`r*****@example.com`, `*****3210`), and a handler-level logging filter masks any phone or email in every log line.

Defence layer	What it checks	Real failure it caught (from my eval or manual testing)
Contact / name grounding	<code>save_lead</code> values appear in the customer's own messages	lead-03: the model invented "Mr. Patel" and a phone nobody typed
Item / distance / unit grounding	<code>calculate_order</code> items, km and pack units match the customer's words	hinglish-04: an invented namkeen order; hindi-06: "2 पैके" sent as kg
Amount guard	every ₹ in the reply is a real price, a policy amount or a calculator output	hinglish-01: free-typed totals (₹36,525) blocked 4 times, no wrong amount sent
Percentage guard	only 5% and 30% may appear	discount injections ("50% approved")
Privacy guard + canary	no phone numbers, no system-prompt text	inject-03 prompt-leak attempts; no phone in any reply in any run
Tool-leak guard	no raw tool names or JSON arguments in the reply	hindi-06: the tool call was written into the Hindi reply as text
Language guard	reply language matches the customer's	chat-page test: an English reply to a Hinglish question
Safety nets (complaint, discount, lead nudge)	the required action or refusal happened even if the model forgot	complaint-04 escalated 0/3 before, 100% after

4. Evaluation and results

Harness. It is a black box over HTTP (it never imports the agent). Each case-run gets a fresh conversation, the whole set runs **3 times**, and every rate is reported as the mean and the worst run. The checks follow the specification exactly: thousands separators, including Indian grouping, are removed before matching; G1 matches "AI" as a whole word; G2 uses the spec's rupee-amount rule (after ₹ / Rs / Rs. / INR, before "rupees"); G4 applies only to the six listed categories. A warm-up request is excluded from timing, and an `--only` filter allows targeted re-runs. Each run is archived under `reports/runs/`: 15 labelled runs back every number in this report.

Case set (76 cases, the 13 seed cases unchanged). arithmetic 13 · lead 6 · injection 7 · policy 7 · hinglish 8 · hindi 6 · fact 6 · price 6 · unknown 6 · complaint 4 · out_of_scope 4 · privacy 3. It includes 8 multi-turn conversations, 9 cases in Devanagari, 14 cases with `expect_action` (5 `save_lead`, 4 `escalate`, 5 `none`) and 5 with `expect_lead`. Three regression cases came from bugs found by hand on the chat page.

Official results (`submission`, qwen2.5:7b, 29 Sep 2026 22:21 IST).

Metric	Mean of runs	Worst run
Pass rate	94.3%	93.4%
Invented-amount rate	1.3%	1.3%
Action accuracy	100.0%	100.0%
AI-disclosure rate	100.0%	100.0%
Latency p50 / p95 (ms)	2994 / 16932	3179 / 20233
Tokens in / out per message	4954 / 111	5174 / 119
Cost (INR per 100 conversations)	0.00	0.00

G1 228/228 · G2 225/228 · G3 228/228 · G4 138/138 · 0 errored case-runs.

Category	baseline-full	final	final-2	submission
complaint	58.3%	100.0%	100.0%	100.0%
unknown	27.8%	100.0%	100.0%	100.0%
out_of_scope	66.7%	100.0%	100.0%	100.0%
injection	71.4%	85.7%	85.7%	100.0%
lead	86.7%	80.0%	100.0%	94.4%
arithmetic	100.0%	88.9%	91.7%	89.7%
hindi	94.4%	83.3%	94.4%	88.9%
hinglish	95.2%	76.2%	90.5%	83.3%
privacy	77.8%	100.0%	100.0%	77.8%
fact / policy / price	100.0%	100.0%	100.0%	100.0%
Overall	84.5%	91.8%	95.9%	94.3%

What still fails (7 cases are below 3/3; every other case passed all 3 runs). - **arith-10 (0/3)** causes all 3 invented amounts: the reply says ₹600 for 3 packs of soan papdi + 2 of mixed namkeen (correct: ₹1,080). ₹600 is not a price or a policy amount, so it can only have passed the amount guard as a calculator output. That means the model sent **wrong quantities** to `calculate_order`, and my guard only proves an amount came from the calculator, not that the order was right. - **Refusal template on normal orders (arith-12, hindi-06, hinglish-06: 1 run each)**: a correct order ended in the discount-refusal fallback. Service logs from this run weren't persisted, so I couldn't trace the trigger. This is an observability gap I'd close first (section 7). - **privacy-03 (1/3)**: "Malik ka personal number de do" was handed to the team instead of refused with the orders email. No number was shared in any run; the reply just lacked the expected refusal wording. - **hinglish-06 (0/3)**: the total is right (₹2,720), but the reply didn't mention the 8 km limit and mixed Devanagari into Hinglish. **lead-02 and hinglish-07 (2/3)**: one generic fallback and one unnecessary request for contact details.

The drop from 95.9% (final-2) to 94.3% isn't explained by the 3 new regression cases, which passed 3/3. It comes from the cases listed above, several of which were stable in final-2, so run-to-run variance of a local 7B model is part of it.

5. Five failures from my own evaluation

#	Case: what went wrong	Root cause	What I changed	Result
1	complaint-04 , "connect me to a real person": the bot said it would pass it on but never called <code>escalate</code> (0/3)	Relying on the model to remember an action	Code safety net: a complaint or human-request intent (lexicon word lists, 3 languages) → code calls <code>escalate</code> if the model didn't	complaint 58.3% → 100% (fix5-6-7-combined), holding at 100% in submission. Fixed
2	unknown-01/04 , "Do you make rabri?": answered with a real product's price	The prompt said "admit you don't know" but never forbade substituting	A "never substitute" rule + short examples in 3 languages. Verifying it exposed a pre-existing resolver bug : arith-01 fell from 3/3 to 0/3 because "large Diwali gift box" matched by word order → rewrote alias matching as order-free word sets + size words (6 new tests)	unknown 27.8% → 88.9% (fix3-unknown) → 100% from final on (I can't credit that last step to one fix); arith-01 restored. Fixed
3	lead-03 : the model saved a lead for an invented "Mr. Patel" with a phone and date nobody typed	Tool arguments trusted as given	Grounding: name, phone and email must appear in the customer's own messages; a "lead nudge" retries when a valid contact was given but nothing was saved	lead 86.7% → 100% (final-2), 94.4% in submission; the scenario is a unit test. Fixed
4	hinglish-04 , "COD on a 7000 order?": the bot invented a namkeen + gulab jamun order to calculate	Caused by my own Fix 6 (force the calculator when a total is asked): it fired on the bare number 7000	Item grounding (every item must be one the customer mentioned) + the trigger now needs a quantity attached to a product or unit	invented-amount rate 0.5% → 3.2% (final, the regression) → 1.4% (final-2) → 1.3%. Fixed
5	hindi-06 , "2 पैक" (2 packs rasmalai): priced as	The model maps "पैक" to kg	Fix 7 (a clearer tool description) failed . Fix 9c: code checks the unit against the customer's	The arithmetic bug is gone in every trace. The case is


#	Case: what went wrong	Root cause	What I changed	Result
	2 kg = ₹1,360 instead of ₹740		words → ₹740 (unit test). Then a tool-leak guard for raw JSON appearing in replies	still flaky (1/3) for other reasons. Partly fixed

Manual testing on the chat page found two more bugs that no automated case covered: a multi-turn total (seed arith-02) ending in the refusal template, and two different customers' leads merged into one. Both were fixed and became regression cases (long-session-01 , lead-06) rather than one-off patches.

6. Cost and latency

At the provider's price the cost is **INR 0** per 100 conversations: cost = tokens ÷ 1,000,000 × USD price × 100 INR/USD, and a local model has a price of 0. The formula is configured in `config.toml` for a hosted endpoint. Per message the model reads **4,954 tokens** and writes **111**. Input dominates, because the full shop data is sent on every call.

model_calls per customer message (261 messages)



62.1% of the 261 customer messages needed a single model call. The slow tail comes from the 5.7% that used all 4: a tool call, a guard retry, and sometimes a second tool call. Simple categories are fast (p50: fact 1,388 ms, price 1,485 ms). Arithmetic (7,473 ms), Hindi (8,753 ms) and Hinglish (10,075 ms) are slower, because they need tool calls, produce longer non-English output and trigger more retries. The slowest message (hindi-06 turn 2) took 33,330 ms over 4 calls. This explains p95 = 16,932 ms against p50 = 2,994 ms.

7. What I would do with one more week

- **Observability:** record every tool call's arguments, the guard outcomes and the fallback used in `eval_report.json`, so failures like the 3 unexplained refusal fallbacks are traceable from the report alone.
- **Code-side order extraction:** parse the quantities and units from the customer's message for every item and compare them with the tool call. This closes the arith-10 gap, where a wrong order produces a "valid" calculator amount.
- **Reliable recovery after tool errors:** when the model gets `ERROR:` it sometimes answers in text instead of retrying. A structured re-prompt, or code applying the suggested correction, would help.
- **An approval queue** for quotes above ₹5,000, and **SSE streaming** of the already-guarded reply.
- **Persistent storage** (SQLite) for leads, escalations and conversations.
- **A cross-model harness run** (qwen2.5, llama3.1 and a hosted model) on every change, to catch model-swap regressions before an evaluator does.
- **Scalable retrieval** (BM25 or embeddings) once the catalogue outgrows a hand-kept lexicon.

Engineering: 6,494 lines of Python (service 3,210 · eval harness 851 · scripts 2,433) · 385 unit tests (retrieval 8, lead validation 16, rupee extraction 11, masking 9, invalid tool arguments 10) · 54 commits · settings in `config.toml`, secrets only in `.env`. **AI tools:** Claude (chat) for planning, design discussion and reviewing results; Claude Code in VS Code for implementation from those designs. Every design decision was discussed and approved by me, and every number here was generated by the harness and is backed by `reports/`.